


Hadoop MapReduce Framework

Overview A MapReduce *job* splits the input dataset into independent chunks processed by *map* tasks in completely parallel manner, followed by *reduce* tasks which aggregate their output. Typically both the input and the output of the job are stored in a Hadoop FileSystem.

MapReduce inputs and outputs The MapReduce framework operates exclusively on $\langle \text{key}, \text{value} \rangle$ pairs. The *keys* and *values* have to be serializable as *Writable*s and additionally the *keys* have to be *WritableComparables* in order to facilitate grouping by the framework.

The following is an overall flow of a MapReduce Job:



Hadoop mapper The mapper takes input $\langle \text{key}, \text{value} \rangle$ pairs, perform computation, and generates a set of intermediate $\langle \text{key}, \text{value} \rangle$ pairs.

```
map(KEYIN key, Iterable<VALUEIN> values,
    Mapper.Context ctx);

public class WordCount {
    public static class Map
        extends Mapper<LongWritable, Text,
            Text, IntWritable> {
        public void map(LongWritable key,
            Text value,
            Context context) {
            String line = value.toString();
            StringTokenizer tk = new ... (line);
            while (tk.hasMoreTokens()) {
                word.set(tk.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

Hadoop reducer The Hadoop reducer reduces a set of intermediate values *which share a key* to a (usually smaller) set of values.

A reducer has three primary phases: *shuffle*, *sort*, and *reduce*. In the **shuffle** phase, input to the Reducer is sorted output of the mappers. In this phase, the framework fetches the relevant partition of the output of all the mappers, via HTTP. In the **sort** phase, the framework groups Reducer inputs by keys in this phase (since different mappers may have output the same key). At the **reduce** phase, the `reduce()` method is called for each $\langle \text{key}, \text{list of values} \rangle$ pair in the grouped inputs.

```
reduce(KEYIN key, Iterable<VALUEIN> values,
    Reducer.Context ctx);

public class WordCount {
```

```
    public static class Reduce
        extends Reducer<Text, IntWritable,
            Text, IntWritable> {
        public void map(Text key,
            Iterable<IntWritable> values,
            Context context) {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }
}
```

The heart of Reducer is its `reduce()` method. This method is called *once per key*. The second argument is an *Iterable* which returns all the values associated with the key.

Running MapReduce jobs

```
public class WordCount {
    public int run(String[] args) {
        Job job = new Job(getConf());
        job.setJarByClass(WordCount.class);

        job.setOutKeyClass(Text.class);
        job.setOutValueClass(IntWritable.class);

        job.setMapperClass(Map.class);
        job.setCombinerClass(Reduce.class);
        job.setReducerClass(Reduce.class);

        job.setInputFormatClass(TextInputFormat);
        job.setOutputFormatClass(TextOutputFormat);
    }

    public static void main(String[] args) {
        ToolRunner.run(new WordCount(), args);
    }
}
```

How Hadoop runs a MapReduce job For a single MapReduce job, there four independent entities: *client*, *job-tracker*, *tasktracker*, and *HDFS*.